Alex Morgan
alex.morgan@email.example | (555) 234-5678 | San Francisco, CA | github.com/alexmorgan-ai

PROFESSIONAL SUMMARY
Results-driven AI/ML Engineer with 3+ years of experience designing and deploying machine learning pipelines, deep learning

TECHNICAL SKILLS
- Languages: Python, SQL, C++, Bash
- ML & AI Frameworks: PyTorch, TensorFlow, Scikit-Learn, Hugging Face, LangChain
- Backend & Cloud: FastAPI, Docker, REST APIs, PostgreSQL, Redis, AWS (S3, EC2), Git
- Core Domains: Natural Language Processing, Computer Vision, MLOps, CI/CD

PROFESSIONAL EXPERIENCE
Machine Learning Engineer | Apex AI Systems, San Francisco, CA (2022 - Present)
- Engineered scalable ML inference pipelines utilizing Python, PyTorch, and FastAPI, serving 1.2M daily predictions with sub-8
- Containerized model training workloads with Docker and orchestrated continuous deployment using GitHub Actions.
- Optimized vector search queries and PostgreSQL data pipelines, improving batch ingestion speeds by 40%.

Associate Software Engineer | DataPulse Labs, San Jose, CA (2021 - 2022)
- Collaborated with engineering teams to develop automated data validation and feature engineering microservices.
- Implemented real-time anomaly detection pipelines using Scikit-Learn and Pandas.

EDUCATION
B.S. in Computer Science | University of California, Berkeley (2017 - 2021)
- Dean's Honors List, Relevant Coursework: Deep Learning, Distributed Systems, Algorithms.

PROJECTS
Autonomous Multi-Agent Workflow Engine
- Developed an open-source agentic workflow framework using Python, FastAPI, and LangChain for automated code generati
Neural Document Retrieval System
- Implemented semantic search engine using PyTorch and FAISS, achieving 94% retrieval accuracy across 500K documents.